

Ramon Moreira Filho

Machine Learning Engineer · MLOps · Data Scientist

Available for international relocation

Belo Horizonte, MG, Brazil · +55 37 99961-0179 · ramon@bring-data.com · linkedin.com/in/ramon-moreira-filho

Languages: Portuguese (native) · English (fluent) · Spanish (fluent)

SUMMARY

Machine Learning Engineer with 7 years in digital lead acquisition (+120 launch cycles) and 4 years focused on end-to-end data science projects. Founder of **Bring Data**, where I single-handedly built a production lead scoring system that sends propensity signals to Meta via the Conversions API in under 5 minutes.

Verified results in real operation: **R\$ 470k of incremental margin in 4 months, +R\$ 0.92 returned per R\$ 1 invested in ads**, higher ROAS in **12 of 12 launches** vs Control (**p<0.05 in 5 of 7 A/B tests**), **CPL 27–44% lower**.

Specialist in end-to-end MLOps, multi-tenant architecture, and production patterns: training-serving parity audit, feature registry, gradual canary deploy, fail-loud asserts, pre-encoding validator.

PROFESSIONAL EXPERIENCE

Machine Learning Engineer — Bring Data

Jan/2025 – Present

Founder and sole developer. ML lead scoring system in production on GCP, multi-tenant architecture.

- **End-to-end system** with RandomForest over 59 features (AUC 0.745); multi-tenant architecture via ClientConfig (153 hardcodes refactored into 15 typed sub-configs in Python) — new clients onboarded with no code changes.
- **Verified results (DevClub, 4 months):** R\$ 470k incremental margin · +R\$ 0.92 returned per R\$ 1 spent on ads · higher ROAS in **12/12 launches** vs Control (**p<0.05 in 5/7 A/B tests**) · **CPL 27–44% lower** · zero cannibalization (Pearson -0.23).
- **Production MLOps:** SSoT src/core/ (15 modules) eliminates training-serving skew by construction; column-by-column **parity audit** between training and production; **feature registry** synchronized via MLflow; mandatory **fail-loud asserts** on critical transforms; gradual **canary deploy** (--no-traffic → 5% → 10% → 100%) with automated **progression gate**; **pre-encoding validator** with offline snapshot.
- **GCP infrastructure:** 31 FastAPI endpoints on Cloud Run + Cloud Scheduler + MLflow (Cloud SQL Postgres + GCS) + Railway PostgreSQL (~100k leads in production). 11 Tier 1 safeguard items delivered incrementally.
- **Monthly retraining** with hook injection (zero training-logic duplication) and quality gate on AUC, monotonicity and positives concentration before promotion. **A/B champion/challenger** in production via UTM routing with distinct CAPI event names — per-model ROAS read directly from Meta Ads Manager.
- **Daily automated monitoring** (feature drift, CAPI/FBP/FBC quality, Meta rejection rate, 12 verification points per cycle) with Slack alerts. **External integrations:** Meta Conversions API, Guru, Hotmart, Google Sheets via Apps Script.

Data Scientist — Data Mundo

Jun/2024 – Jan/2025

- End-to-end projects: EDA, predictive modeling, production delivery.
- ETL pipelines with SQL, Pandas and Apache Spark for large data volumes.
- Predictive models (Random Forest, XGBoost, LightGBM) for churn and campaign optimization — **+20% ROI**.
- Power BI and Tableau dashboards; executive reports for non-technical stakeholders.

Founder & Data Director — Produções de Los Angeles

Jan/2017 – Jun/2024

- Data-driven strategies for digital campaigns — over **US\$ 450k annually** in revenue.
- Lead scoring with supervised ML — **15% reduction in lead acquisition cost**.
- Real-time data collection integrations via API (Google Ads, Meta Ads).
- **+120 launch cycles** in digital education products — deep domain knowledge of acquisition funnels.

MAIN PROJECT — BRING DATA

ML system that classifies leads into deciles D1–D10 and sends a calibrated signal to Meta via Conversions API in <5 min — Meta's algorithm optimizes for profiles with higher real conversion probability, not just form completions.	
Architecture	FastAPI · Cloud Run · Cloud Scheduler · MLflow (Cloud SQL + GCS) · Railway PostgreSQL · scikit-learn (RandomForest, 59 features) · facebook-business · Apps Script
MLOps patterns	Parity audit · feature registry · fail-loud · canary deploy · pre-encoding validator · per-platform CAPI dispatch · A/B via UTM with distinct event names
Outcome	R\$ 470k incremental · +R\$ 0.92 / R\$ 1 · 12/12 launches · p<0.05 in 5/7 A/B · CPL –27% to –44% · top-2 deciles convert 3× above average

TECHNICAL SKILLS

Languages	Python, SQL, R, JavaScript (Node.js)
ML	Scikit-learn, XGBoost, LightGBM, RandomForest, SHAP, TensorFlow, PyTorch
MLOps Patterns	Training-serving parity audit, feature registry, fail-loud asserts, gradual canary deploy, pre-training schema validation, UTM-routed A/B, per-platform dispatch, retraining with quality gate
MLOps Tooling	MLflow, FastAPI, Pydantic, SQLAlchemy, Docker, Cloud Run, Cloud Scheduler, Git/GitHub, CI/CD
Cloud / Infra	GCP (Cloud Run, Cloud SQL, Cloud Storage, BigQuery, Cloud Logging) · AWS (S3, EC2, SageMaker)
Databases	PostgreSQL (Railway, Cloud SQL), MySQL, BigQuery, SQLite
Visualization	Power BI, Tableau, Plotly, Matplotlib, ggplot2
Big Data	Apache Spark, Apache Airflow
Integrations	Meta Conversions API (facebook-business), Meta Ads API, Google Ads API, Google Sheets API + Apps Script, Guru API, Hotmart API

EDUCATION & CERTIFICATIONS

Stanford / DeepLearning.AI	Machine Learning Specialization — Andrew Ng
DeepLearning.AI	Machine Learning in Production — MLOps Specialization
Google	Machine Learning Engineer Professional Certificate
University of Michigan	Python for Everybody
IBM / Coursera	Databases and SQL for Data Science
MLOps Community	First Stack MLOps
Data Mundo	Data Science Certification